The small refit-versus-odds gap changes across splits

CONTROLLED PRIOR-SHIFT SCENARIOS | q = 1% | paired test log-loss difference

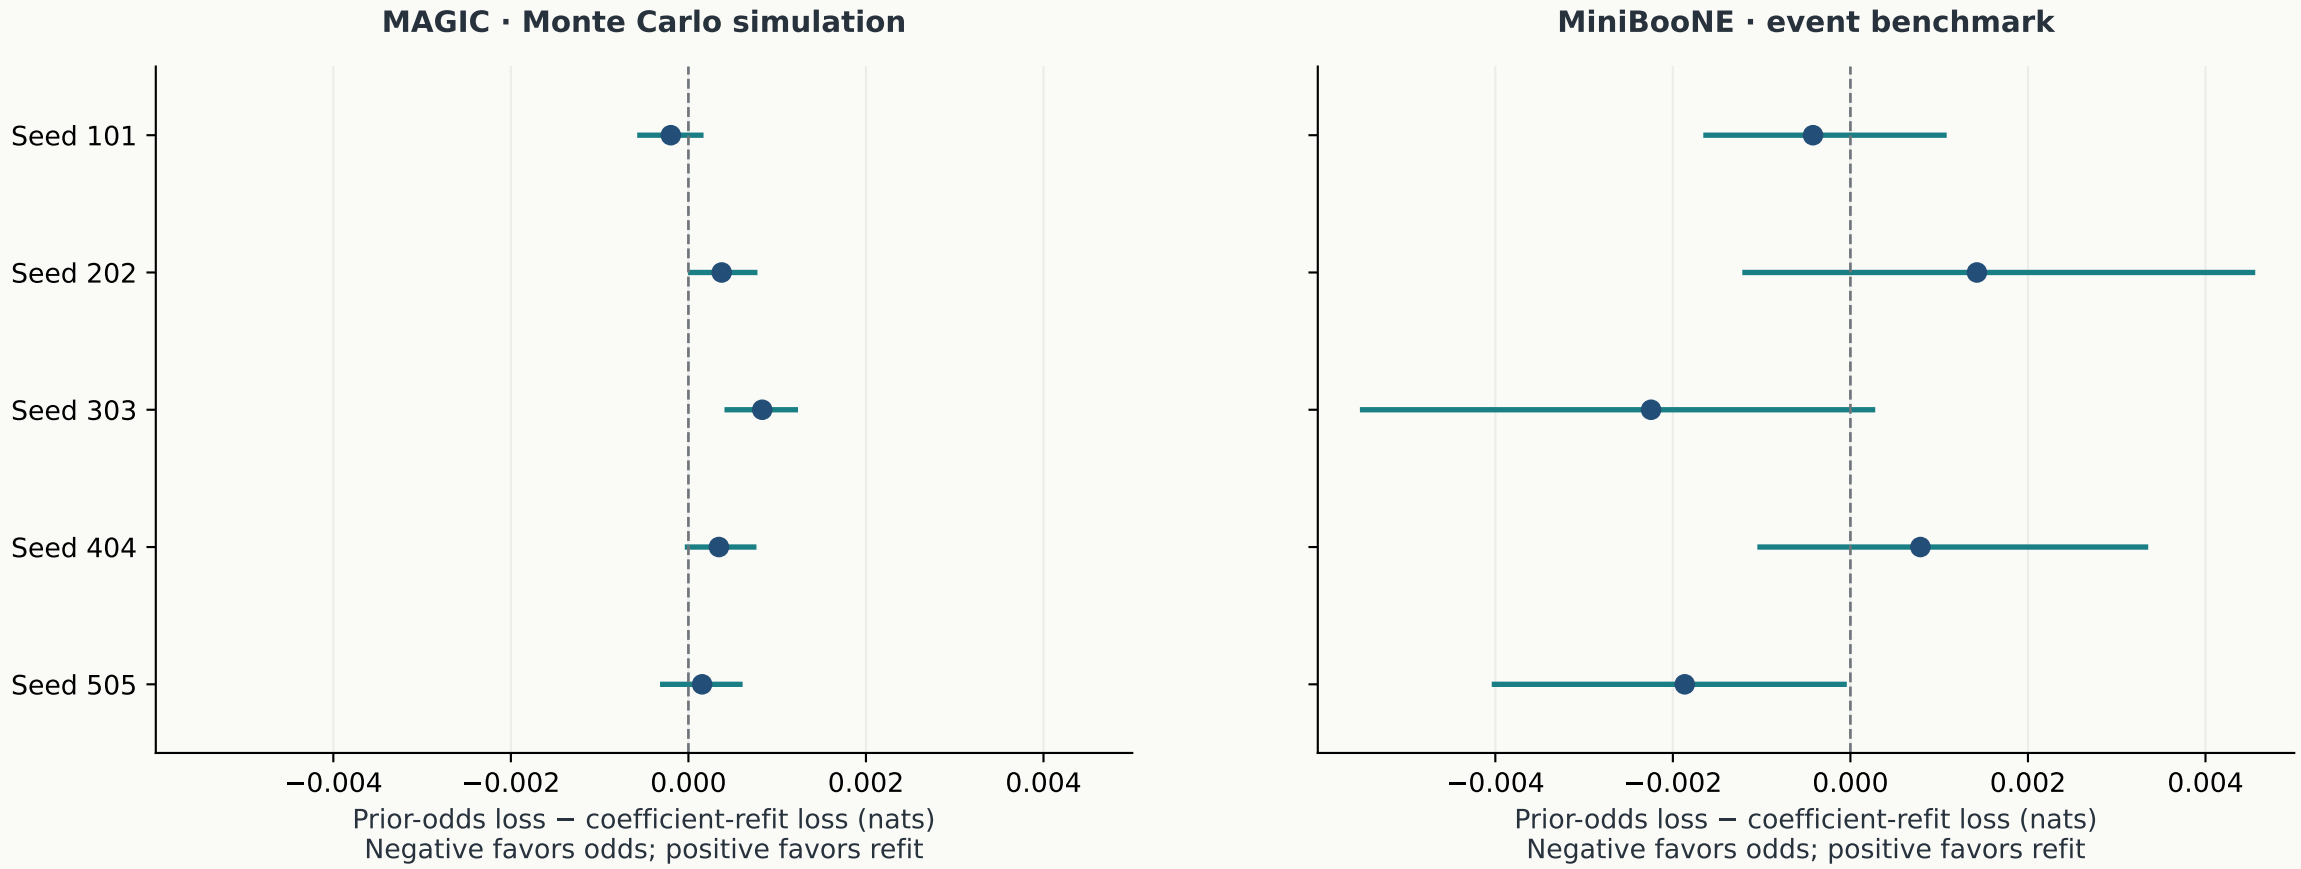

Lines: conditional 95% percentile intervals from 1,000 whole-feature-group draws per split. No pooled significance test.
Split seeds overlap. Intervals assume these benchmark groups and the supplied target prevalence; unknown run dependence remains.